

01

The Playbook

Why your AI videos fall apart after the third shot — and the production system that fixes it for good.

Start here

If you have ever generated nine shots for one video and had the face change in four of them, this book is the fix. Not a better prompt. A different way of working.

Here is what almost everybody does. They open an AI video tool, type a prompt, and hit generate. They get a clip. It is 70% right. So they change three words and generate again. And again. Twenty generations later they have spent real money, they have a folder full of near-misses, and the character's face is different in every single one.

Then they do it eight more times for the other shots in the video. By the time they stitch it together, it does not look like one film. It looks like eight different films about eight different people.

Most people conclude the models are not good enough yet. **They are wrong.** The models are extraordinary. The *workflow* is what is broken.

THE ONE IDEA IN THIS ENTIRE COURSE

A video is not one generation. It is a graph of shots — and every shot needs to be something you can regenerate on its own, without touching the others, anchored to an identity you locked down once.

That is it. Everything else in these three books is the practical consequence of taking that sentence seriously. If you only change one thing about how you work, change that.

What you will be able to do when you finish

- Put the same recognizable person in twelve different shots and have them read as one person.
- Fix a single bad shot in a finished video without regenerating anything else.
- Know which model to use for which shot instead of forcing one model to do everything.
- Cut your generation spend by more than half, mostly by stopping the waste.
- Hand a finished piece to an editor as a real timeline instead of a locked MP4.
- Turn one generated video into five published posts without generating anything new.

Why the slot machine costs you everything

Consumer AI video apps are built around a single interaction: prompt in, clip out. That interaction is fine when you want one clip. It is catastrophic when you want a film, and here is precisely why.

WHAT HAPPENS	WHY IT COSTS YOU
Shot 7 of 9 comes out wrong	No way to fix just that shot. You regenerate the batch or you live with it.
A take comes out great	No record of exactly what produced it. You cannot reproduce it tomorrow.
You generate the same character twice	The model has no memory. It invents a slightly different face each time.
You finish the video	You get a flat MP4. Captions are baked in. You cannot recut it, recrop it, or change the hook.
You want a second version	You pay full price again for something you already generated.

Notice that four of those five are not quality problems. They are *bookkeeping* problems. The tool threw away information you needed. That is why the answer is a system, not a subscription.

THE EXPENSIVE MISTAKE

Regenerating a whole video because one shot is wrong is the single biggest money leak in AI video. If a 12-shot video costs \$50 to generate and you do that three times to fix three different shots, you spent \$150 to deliver \$50 of work. The people who make this profitable are not getting better clips. They are re-rolling one shot at a time.

The system, end to end

Four layers. Each one solves a specific failure above. Build them in this order.

LAYER 1 — IDENTITY

A locked set of reference images and a voice, built once, referenced by every shot forever. This is what stops the face drifting. It is the highest-leverage thing you will build.

LAYER 2 — THE SHOT GRAPH

Your video written as a list of individually addressable shots, each with its own prompt, its own model, and its own version history. This is the project file the apps never gave you.

LAYER 3 — ROUTING & TAKES

Each shot goes to the model that is actually best at that *kind* of shot, and you generate three candidates and pick one. Choice is the product you are buying.

LAYER 4 — ASSEMBLY

Captions, music, titles and logos are added at the *end*, never baked into generation. This is what makes five variants free.

Layer 1 — Lock the identity

Everything downstream depends on this and most people skip it. They grab three photos off a phone, feed them in, and wonder why the jaw changes between shots.

What a real identity kit contains

Shoot this deliberately. It takes 45 minutes and you will use it for a year. A controlled phone shoot beats a folder of existing photos every time, because existing photos have inconsistent lighting, heavy retouching, and usually only one or two angles.

GROUP	WHAT TO CAPTURE	WHY
Angles	Front, 3/4 left, 3/4 right, profile left, profile right, slight low, slight high	A profile shot built only from front-facing refs will invent a profile. You need the real one.
Expressions	Neutral, smile, laugh, stern, mid-speech, skeptical, thinking	Emotion changes face shape. Give the model the truth instead of letting it guess.
Framing	Close head, chest up, waist up, full body standing	Wide shots need body proportion references or you get a stranger at distance.
Wardrobe	The full angle set again, in each signature outfit	Mismatched wardrobe refs are why the jacket changes between cuts.
Voice	3 minutes clean: neutral, excited, serious, plus unscripted	30 seconds gets you a clone. Three minutes gets you a good one.

SHOOTING CONDITIONS THAT MATTER

- **Flat, soft, even light.** Window light or one softbox. No hard shadows, no color casts.
- **Plain background.** A blank wall. Busy backgrounds bleed into the references.
- **No filters, no beauty mode, no portrait blur.** Every one of those removes information the model needs.
- **Highest resolution the phone offers.** You can always downscale.
- **Audio in a quiet room.** No HVAC, no traffic, no echo. A closet with clothes in it is a genuinely great vocal booth.

Target 40 to 60 usable frames. You will not use all of them at once — you will select three to five per shot. More references is not better. *Better matched* references is better.

THE RULE THAT SAVES YOU THE MOST RE-ROLLS

Match the reference angle to the shot angle. Generating a profile turn? Pass the profile references. Generating a low-angle hero shot? Pass the low-angle reference. Most face drift people blame on the model is actually them passing front-facing refs for a shot that is not front-facing.

The six rules that keep a face locked

- 1 **Never mix identity versions inside one video.** If you rebuild your kit with better photos, rebuild the whole video against it or none of it. Drift between shots is the most obvious tell there is.
- 2 **Match reference angle to shot angle.** See above. This is the big one.
- 3 **Match wardrobe references to scene wardrobe.** Otherwise the model splits the difference and invents a third outfit.
- 4 **Never route a talking-to-camera shot through text-to-video.** Use an audio-driven tool. Generated mouths on a close-up are the fastest way to look cheap. More on this in Layer 3.
- 5 **Once a take is approved, never regenerate it.** Even at the same settings you will get a subtly different face. Locked is locked. "Just one more to see" has cost people entire days.

6

Keep a notes file. Log which model drifted on which shot type. After two weeks you will have a routing playbook nobody else has.

Layer 2 — Write the shot graph

Before you generate anything, your video exists as a written list of shots. This is the single habit that separates people who ship from people who burn credits.

Every shot carries: an ID, a duration, what kind of shot it is, what happens in frame, the camera framing, any dialogue, which identity references to use, and — once you start generating — a log of every take and which one you approved.

ID	TYPE	LEN	ACTION	REFS
s01	story	3s	She slides a check across the desk, grinning	—
s02	broll	4s	Insert: pen circles the monthly payment	—
s05	story	5s	He speaks from the waiting chairs behind her	3/4 left, skeptical, chest up
s11	direct	9s	To camera, calm, delivers the point	front, speaking, chest up

You can keep this in a spreadsheet, a Notion table, or a text file. The format does not matter. **What matters is that each row is a thing you can regenerate by itself.**

Why the shot graph pays for itself immediately

- You catch problems before you spend. Reading twelve rows takes two minutes and reveals the pacing problem you would otherwise pay to discover.
- You can price the job accurately before starting.
- When a client says "I don't like the third one," you know exactly what to re-roll.
- Your good prompts become reusable assets instead of scrollback you will never find again.

Naming and filing your takes

Every generated clip gets saved as `shotID.version.mp4` — `s04.v1.mp4` , `s04.v2.mp4` , `s04.v3.mp4` . Nothing is ever overwritten. Nothing is ever deleted. Storage is cheap; regenerating a take you already had is not.

APPROVE, THEN MOVE ON

When you pick a winner, write it down — "s04 → v2." That one note is what lets you assemble the video, hand it to an editor, come back in three weeks and still know which file was the good one. Most people's biggest time sink is re-watching takes they already judged.

Layer 3 — Route each shot to the right model

There is no best AI video model. There are models that win at specific things. Forcing one model to do everything is why your output is uneven and your bill is high.

SHOT TYPE	ROUTE IT TO	WHY
Talking to camera	An audio-driven avatar/lip-sync tool	Real lip-sync from real audio. Never generate a talking close-up from text.
Character in a scene	A multi-reference model (accepts many reference images)	Best identity hold across many separate shots.
Several angles in one clip	A multi-shot model	Keeps identity across cuts inside a single generation. Cheap per second.
Hero / establishing	Your most expensive, highest-fidelity model	This is the shot people screenshot. Spend here.
B-roll, no face	Your cheapest model	Identity does not matter. Do not pay premium rates for a shot of a door.

Book Two has the current model cheat sheet with per-second prices, and it is built to be updated — because it will change, probably within a month of you reading this. The *routing logic* above will not change. Learn the logic, look up the model.

Three candidates, always

Generate three takes of every shot. Not one. Three.

This feels wasteful and it is the opposite. One take means you either accept what you got or pay again. Three takes means you *choose*. The cost difference is small; the quality difference is the entire product. Professional photographers do not take one frame. Neither should you.

WHERE PEOPLE WASTE MONEY INSTEAD

The waste is not in generating three candidates. The waste is generating one candidate twelve times because you keep tweaking the prompt. Write the prompt carefully once, generate three, pick one, move on. If all three are wrong, the prompt is wrong — fix the prompt, do not spin.

Layer 4 — Assemble last, and keep it editable

This is the layer that turns one video into five, and it depends entirely on one discipline:

NEVER BAKE ANYTHING INTO GENERATION

No captions. No titles. No logo. No music. No end card. Generate clean footage only, and add every one of those at the edit. If a caption is burned into a generated clip, changing it means regenerating the clip. If it is a layer in your editor, changing it takes eight seconds.

The variant engine

Once you have a finished piece assembled from clean footage, you own five products, not one:

VARIANT	WHAT CHANGES	WHERE IT GOES
A — Story cut	The full piece as intended	TikTok, Reels, Shorts
B — Hook swap	Different first 3 seconds, same body	Repost 6+ days later
C — Short cut	Just the payoff, ~12s, captioned	Reels, LinkedIn
D — Landscape	Recrop to 16:9, more breathing room	YouTube, X
E — Still	Best frame + the key line as a graphic	Instagram, LinkedIn

Generation is the expensive part. Assembly is basically free. If you are generating a new video for every post, you are paying five times for something you already own.

Keep the edit alive

Whatever editor you use, export a project file, not just an MP4. When a client comes back in six weeks wanting the CTA changed, the difference between "sure, an hour" and "we'd have to rebuild it" is whether you kept the timeline. Book Three treats this as a business asset, because it is one.

The production loop

Once the identity kit exists, every video runs the same seven steps. It should get boring. Boring is the goal — boring is what you can do twice a day.

- 1 **Write the shot graph.** Twelve rows. Do not open a generation tool yet.
- 2 **Price it.** Sum the shot durations, multiply by candidates, multiply by rate. Know the bill before you spend.
- 3 **Generate.** Three candidates per shot, all shots, one pass.

- 4 **Review and approve.** Watch every take. Pick one per shot. Write it down.
- 5 **Re-roll the failures.** Only the shots that failed. Never the whole video.
- 6 **Assemble.** Clean footage in, then captions, music, titles, end card.
- 7 **Cut the variants.** Five outputs from one generation.

The review pass that protects your reputation

Watch the assembled cut twice. Once at full speed with sound — does the opening hold you, does it land, does the lip-sync hold on the close-ups. Then **once muted at double speed**. Face drift and hand artifacts leap out when you are not listening to the words. This second pass takes ninety seconds and catches most of what the first one misses.

Kill criteria

Do not publish if any of these are true. Re-roll instead.

- ☐ The face differs noticeably between any two shots.
- ☐ A hand does something anatomically impossible in a shot the eye lands on.
- ☐ Lip-sync drifts on a close-up.
- ☐ Text inside the generated frame is garbled. (It always will be. Add text at the edit.)
- ☐ It feels like a slideshow instead of a scene.

ONE GOOD VIDEO BEATS TWO MEDIOCRE ONES

Every platform learns what quality level your account produces and distributes accordingly. Publishing something you know is soft to hit a quota trains the algorithm against you, and that is expensive to undo. When in doubt, ship one.

What to do this week

- 1 **Build one identity kit.** Yourself, a friend, or a generated character — it does not matter. 45 minutes. This unblocks everything.
- 2 **Write one shot graph.** Twelve rows, 45 to 60 seconds. Do not generate yet.
- 3 **Generate three candidates for one single shot.** Just one shot. Look at the three side by side. That comparison teaches you more than twenty solo generations.
- 4 **Then run the full loop once,** start to finish, on that one video. Book Two has every prompt and template you need. The videos walk you through both of these live.

Next: **Book Two — The Vault.** The prompt formulas, shot templates, model cheat sheet and troubleshooting table. That is the one you will keep open on a second monitor while you work.