

02

The Vault

Every prompt formula, shot template, model comparison and troubleshooting fix. Built to be kept open on a second monitor.

How to use this book

Book One taught the system. This is the parts bin. Nothing here needs to be read in order — find your shot type, copy the template, fill the slots, generate.

- 1 The house prompt formula
- 2 Shot templates — 12 you will use constantly
- 3 Camera and lighting language that actually works
- 4 The identity kit shot list (print this)
- 5 Model cheat sheet and what things cost
- 6 Troubleshooting — symptom, cause, fix
- 7 The shot graph template

1 · The house prompt formula

Almost every bad AI video prompt fails the same way: it describes a *vibe* instead of a *shot*. Models are very good at rendering a specific instant and very bad at interpreting a mood. So we write in slots, always in this order.

```
[SUBJECT]. [ACTION]. [CAMERA]. [SETTING + LIGHT]. [REALISM BLOCK]. [NEGATIVE BLOCK].
```

SLOT	WHAT GOES IN IT	COMMON MISTAKE
Subject	Who is on screen, tied to your references. Wardrobe if it matters.	Re-describing the face. Let the references do it.
Action	One continuous physical action. Present tense.	Three actions in sequence. A 5-second clip does one thing.
Camera	Framing plus one movement.	Naming four movements at once.
Setting + light	Where, what time, what the light is doing.	Leaving it out and getting a grey void.
Realism block	Fixed text. Never changes. See below.	Skipping it and getting plastic skin.
Negative block	Fixed text. Never changes. See below.	Skipping it and getting seven fingers.

THE REALISM BLOCK — PASTE AT THE END OF EVERY PROMPT

shot on a cinema camera, natural skin texture with visible pores, practical lighting, subtle handheld motion, shallow depth of field, filmic color grade, no beauty retouching

THE NEGATIVE BLOCK — PASTE AFTER IT

no plastic skin, no waxy face, no warped hands, no extra fingers, no morphing facial features, no floating text, no watermark, no distorted faces in background, no sudden identity change

WHY "NO BEAUTY RETOUCHING" MATTERS MORE THAN IT SOUNDS

Models are trained on enormous amounts of retouched imagery, so their default human is airbrushed and slightly plastic — and that plastic look is the number one reason a viewer clocks a video as AI in under a second. Explicitly asking for pores and natural texture is the cheapest quality upgrade in this entire course.

A FILLED-IN EXAMPLE

A composed man in his fifties in a navy suit, matching the reference images, sits in a waiting-room chair with a folder on his knee. He looks up from the folder and speaks one calm sentence. Medium shot, slow push in, eye level. Car dealership finance office, late afternoon light through venetian blinds, warm practical lamps. Shot on a cinema camera, natural skin texture with visible pores, practical lighting, subtle handheld motion, shallow depth of field, filmic color grade, no beauty retouching. No plastic skin, no waxy face, no warped hands, no extra fingers, no morphing facial features, no floating text, no watermark, no distorted faces in background, no sudden identity change.

2 · Twelve shot templates

Fill the brackets. Append the realism and negative blocks to every one of them — they are omitted below only to save space.

Dialogue and performance

DIRECT ADDRESS (ROUTE TO AUDIO-DRIVEN LIP-SYNC, NEVER TEXT-TO-VIDEO)

[SUBJECT] speaks directly to camera, calm and grounded, minimal hand movement. Medium close-up, eye level, locked off. [SETTING], soft key light from camera left, subtle falloff behind.

OVERHEARD LINE (SUBJECT SPEAKS TO ANOTHER CHARACTER)

[SUBJECT] turns slightly and says one short line to someone just off frame, then returns to what they were doing. Medium shot, slow push in, eye level. [SETTING], [LIGHT].

REACTION, NO DIALOGUE

[SUBJECT] listens, and their expression shifts from [EMOTION A] to [EMOTION B]. No speech. Close-up, static, shallow focus. [SETTING], [LIGHT].

Movement and blocking

WALK-IN ENTRANCE

[SUBJECT] walks into frame from the left, stops, and looks toward [TARGET]. Wide shot, camera static at chest height. [SETTING], [LIGHT].

SEATED AUTHORITY

[SUBJECT] sits back in a chair, hands resting still, holding a steady gaze. Medium shot, very slow push in, slightly low angle. [SETTING], [LIGHT].

PROFILE TURN (PASS YOUR PROFILE REFERENCES)

[SUBJECT] stands in profile and turns to face camera. Medium close-up, static, eye level. [SETTING], [LIGHT].

Coverage and inserts

OBJECT INSERT

Extreme close-up of [OBJECT] on [SURFACE]. A hand enters frame and [SIMPLE ACTION]. Macro lens, shallow focus, static. [LIGHT].

ESTABLISHING EXTERIOR

Wide establishing shot of [PLACE] at [TIME OF DAY]. No people in foreground. Slow drift right. [WEATHER AND LIGHT].

ENVIRONMENTAL B-ROLL

[DETAIL OF SETTING] with slow ambient motion – [dust in light / traffic behind glass / curtain moving]. Static, shallow focus. [LIGHT].

Story beats

THE TENSION BEAT

[SUBJECT A] and [SUBJECT B] face each other across [OBJECT]. Neither speaks. Medium two-shot, static, eye level. [SETTING], hard directional light creating separation.

THE REALIZATION

[SUBJECT] reads [OBJECT] and stops moving. The expression settles. Close-up, slow push in, shallow focus. [SETTING], [LIGHT].

THE EXIT

[SUBJECT] walks away from camera down [PATH] and does not look back. Wide shot, static, low angle. [SETTING], [TIME OF DAY] light.

3 · Language that actually works

Models respond to real film vocabulary because they were trained on real film. Vague adjectives get you vague output. Use these words.

FRAMING — USE EXACTLY ONE

extreme close-up · close-up · medium close-up · medium shot · medium wide · wide shot · extreme wide · over-the-shoulder · two-shot · point of view · insert · macro

MOVEMENT — USE EXACTLY ONE

static / locked off · slow push in · slow pull out · pan left / right · tilt up / down · tracking alongside · handheld follow · slow drift · rack focus · crane up

ANGLE

eye level · slightly low angle · low angle · slightly high angle · high angle · overhead · dutch angle (use sparingly)

LIGHT — THE MOST UNDERUSED SLOT

soft key from camera left · hard directional sunlight · overcast diffuse · golden hour backlight · practical lamps only · window light with visible falloff · single source with deep shadow · fluorescent overhead · firelight flicker · neon spill

ONE MOVEMENT PER SHOT

"Slow push in while panning right and tilting up" produces a mess in almost every model. Real cinematography usually does one thing per shot too. Pick one. If you need two moves, that is two shots — which is fine, because in this system shots are cheap and separable.

WORDS THAT QUIETLY HURT YOU

AVOID	USE INSTEAD
"cinematic"	Name the actual lens, light and movement
"beautiful", "stunning", "epic"	Nothing. These do not describe a frame.
"4K", "8K", "HD"	Set resolution in the tool, not the prompt
"masterpiece", "award-winning"	Nothing. Leftover habits from early image models.
"he says 'X' while walking and pointing"	One action. Split into two shots.
Text you want on screen	Add it in the editor. Generated text is always garbled.

4 • The identity kit shot list

Print this page. Hand it to whoever is holding the camera. Forty-five minutes, one plain wall, one soft light source, no filters, no portrait mode, highest resolution available.

ANGLES — SHOOT EVERY ONE, NEUTRAL EXPRESSION

- ☐ Front, straight on, eyes to lens
- ☐ Three-quarter left
- ☐ Three-quarter right
- ☐ Full profile left
- ☐ Full profile right
- ☐ Slightly low angle (camera at chest, looking up)
- ☐ Slightly high angle (camera above eyeline)

EXPRESSIONS — FRONT-ON FOR EACH

- ☐ Neutral, relaxed
- ☐ Warm closed-mouth smile
- ☐ Genuine laugh
- ☐ Serious / stern
- ☐ Mid-sentence, mouth open, speaking
- ☐ Skeptical, one eyebrow raised
- ☐ Thinking, looking slightly off-camera

FRAMING

- ☐ Tight head and shoulders
- ☐ Chest up
- ☐ Waist up
- ☐ Full body, standing, arms relaxed

WARDROBE — REPEAT THE SEVEN ANGLES IN EACH OUTFIT

- ☐ Signature outfit one (the "on duty" look)
- ☐ Signature outfit two (dressed down)

VOICE — QUIET ROOM, PHONE OR USB MIC, NO BACKGROUND NOISE

- ☐ A neutral paragraph, read plainly (~45s)
- ☐ An excited paragraph, more energy (~45s)
- ☐ A serious paragraph, slower (~45s)

- ☐ Thirty seconds of unscripted natural talking
- ☐ Save the raw uncompressed file somewhere permanent

IF THE PERSON IS NOT YOU

Get written permission before you shoot, not after. A short signed release covering AI-generated video and voice, the channels it will appear on, how long it lasts, and their right to pull it. Book Three has the full breakdown and a template to adapt. This protects them and it protects you, and having it ready is one of the things that makes you look like a professional rather than a hobbyist.

5 · Model cheat sheet

Prices and model names in this category change every few weeks. **The routing logic does not.** Learn the left column; look up the right column before every project.

JOB TO BE DONE	WHAT TO LOOK FOR	ROUGH RATE
Talking to camera	Audio-driven avatar tool that animates a still from a voice track	Per clip, ~\$0.30–0.60
Character in a scene	Multi-reference video model accepting many reference images	Premium, ~\$0.45–0.50/sec
Multi-angle single clip	Model advertising multi-shot with identity continuity	Cheap, ~\$0.03–0.14/sec
Hero / establishing	Highest-fidelity flagship model, native audio	Top tier, ~\$0.05–0.75/sec by quality
B-roll, no face	Whatever is cheapest that looks acceptable	Budget, ~\$0.03/sec
Voice clone	30s minimum sample, expressive delivery, commercial rights	Per 1k characters, cents

What a video actually costs

Do this arithmetic *before* you generate. It takes two minutes and it is the difference between a profitable project and a surprise.

```
total seconds to generate = (sum of all shot durations) × (candidates per shot)
cost = total seconds × rate per second
```

A worked example — a 52-second video, twelve shots, three candidates each, mixed routing:

BUCKET	FINAL SEC	GENERATED	COST
7 story shots @ premium	35	105	~\$49
3 b-roll @ budget	15	45	~\$4
2 direct address @ per clip	10	—	~\$2
Voice + music	—	—	~\$1
Total	52s	150s	~\$56

Route those story shots to the budget model instead and the same video lands near **\$18**. That is the entire value of routing: same runtime, a third of the bill, and you only spend premium where the viewer is actually looking.

REMEMBER THE MULTIPLIER

That \$56 video becomes five published posts through the variant engine. So your real number is about \$11 per published asset — and that is the number to quote when someone tells you AI video is expensive.

6 · Troubleshooting

Find your symptom. Almost every problem in AI video traces to one of these fifteen causes.

SYMPTOM	ACTUAL CAUSE	FIX
Face changes between shots	Different or mismatched references per shot	Same identity kit, same version, angle-matched refs
Face wrong on a profile shot	Only front-facing refs supplied	Pass the profile references for turning shots
Outfit changes between cuts	Wardrobe refs do not match scene wardrobe	Pass refs from the matching wardrobe folder
Skin looks plastic or waxy	Realism block missing	Append the realism block; add "visible pores"
Hands mangled	Hands doing something complex in frame	Reframe so hands are out of shot or still; re-roll
Mouth looks wrong when speaking	Talking shot generated from text	Route to audio-driven lip-sync. Non-negotiable.
On-screen text is gibberish	You asked the model for text	Never. Add all text in the editor.
Clip feels like a slideshow	No movement specified	Add exactly one camera movement
Camera does something wild	Multiple movements specified	One movement per shot
Subject does not do the action	Too many actions in one prompt	One continuous action. Split into two shots.
Everything looks grey and flat	No lighting specified	Fill the light slot every single time
Background people look melted	Crowd in a shallow scene	Add to negative block; or shoot tighter
Costs running away	Re-rolling whole videos, or prompt-spinning	Re-roll single shots. Three candidates then move on.
Good take, cannot reproduce it	Prompt not saved	Log every prompt with its take. Always.
Client wants a change, whole job re-do	Captions and music baked into generation	Generate clean. Add everything at the edit.

7 · The shot graph template

Copy this into a spreadsheet. One row per shot. Fill it out completely before you generate anything — this is where the money is saved.

ID	TYPE	SEC	ACTION	CAMERA	REFS	TAKES	LOCKED
s01	story	3					
s02	broll	4					
s03	story	4					
s04							
s05							
s06							

COLUMN DEFINITIONS

- **Type** — story, broll, direct, hero, multi-angle. Drives which model you route to.
- **Refs** — which identity references to pass. Angle-matched to the shot.
- **Takes** — filenames generated so far: `s04.v1` , `s04.v2` , `s04.v3` .
- **Locked** — the approved take. Assembly only ever uses these.

Pre-flight checklist

- ☐ Every shot has a type, a duration and one clear action
- ☐ No shot is longer than six seconds except direct address
- ☐ Every shot with the character has angle-matched references listed
- ☐ Wardrobe is consistent across the whole piece
- ☐ No shot asks the model to render text
- ☐ Direct-address shots are routed to lip-sync, not text-to-video
- ☐ You have run the cost math and you are okay with the number

Next: **Book Three** — **The Business**. What to charge, how to run a client, the release template, and how to turn this system into recurring revenue.